

Simple_Regression_Intro

Emmenta Janneh

2024-02-01

We load in the data

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.3.2
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.2      v readr      2.1.4
## v forcats    1.0.0      v stringr   1.5.0
## v ggplot2    3.4.3      v tibble    3.2.1
## v lubridate  1.9.2      v tidyr     1.3.0
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
hibbs <- read.csv("https://dcgerard.github.io/stat_415_615/data/hibbs.csv")
```

You can use `lm()` (linear model) function to find the OLS estimates. - First argument is a “formula” which uses a tilde $y \sim x$ - Second argument is the data frame containing the variables - You assign the output to some variables, so you can manipulate it later

```
lm_hibbs <- lm(vote ~ growth, data = hibbs)
lm_hibbs
```

```
##
## Call:
## lm(formula = vote ~ growth, data = hibbs)
##
## Coefficients:
## (Intercept)      growth
##      46.248         3.061
```

To interact with this `lm` object, it is easiest to use the `broom` package. - `glance()`: will provide a 1 row summary of the model - `tidy()`: will provide one row per parameter estimate - `augment()`: will provide one row per observational unit

```
library(broom)
```

```
## Warning: package 'broom' was built under R version 4.3.2
```

```
tidy(lm_hibbs)
```

```
## # A tibble: 2 x 5
##   term          estimate std.error statistic  p.value
##   <chr>         <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)   46.2      1.62     28.5 8.41e-14
## 2 growth        3.06     0.696     4.40 6.10e- 4
```

```
glance(lm_hibbs)
```

```
## # A tibble: 1 x 12
##   r.squared adj.r.squared sigma statistic p.value    df logLik   AIC   BIC
##   <dbl>      <dbl> <dbl>    <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
## 1    0.580      0.550  3.76     19.3 0.000610     1 -42.8  91.7  94.0
## # i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

```
augment(lm_hibbs)
```

```
## # A tibble: 16 x 8
##   vote growth .fitted .resid .hat .sigma .cooksd .std.resid
##   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl>   <dbl>   <dbl>
## 1  44.6   2.4    53.6 -8.99 0.0711 2.92 0.235    -2.48
## 2  57.8   2.89   55.1  2.67 0.0962 3.83 0.0296     0.746
## 3  49.9   0.85   48.8  1.06 0.100  3.89 0.00491    0.297
## 4  61.3   4.21   59.1  2.21 0.246  3.84 0.0742     0.675
## 5  49.6   3.02   55.5 -5.89 0.106  3.50 0.162    -1.66
## 6  61.8   3.62   57.3  4.46 0.164  3.66 0.165     1.30
## 7  49.0   1.08   49.6 -0.603 0.0854 3.90 0.00131   -0.168
## 8  44.7  -0.39   45.1 -0.354 0.242  3.90 0.00186   -0.108
## 9  59.2   3.86   58.1  1.11 0.194  3.89 0.0130     0.328
## 10 53.9   2.27   53.2  0.745 0.0672 3.90 0.00151     0.205
## 11 46.6   0.38   47.4 -0.861 0.141  3.90 0.00501   -0.247
## 12 54.7   1.04   49.4  5.31 0.0877 3.59 0.105     1.48
## 13 50.3   2.36   53.5 -3.20 0.0698 3.80 0.0292   -0.882
## 14 51.2   1.72   51.5 -0.272 0.0636 3.90 0.000189  -0.0746
## 15 46.3   0.1    46.6 -0.234 0.173  3.90 0.000488  -0.0683
## 16 52     0.95   49.2  2.84 0.0932 3.82 0.0324     0.794
```

Note: **We use `augment()` to get the residuals and fitted values.

Lets take a look at another sample


```
data("mtcars")
lm_mtcars <- lm(mpg ~ wt, data = mtcars)
tidy(lm_mtcars)
```

```
## # A tibble: 2 x 5
##   term      estimate std.error statistic  p.value
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)  37.3      1.88     19.9 8.24e-19
## 2 wt          -5.34     0.559    -9.56 1.29e-10
```

Cars that weight 1000 pounds more have 5.3 worse MPG on average. 37.285 is the y-intercept of the regression line.

You should always plot the data BEFORE doing a regression

```
ggplot(mtcars, aes(x = wt, y = mpg)) +
  geom_point()
```



What if you want an estimated value at an X not in the data set? To predict what Y will be given a value of X , you create a new data frame with those values of X .

```
newdf <- data.frame(growth = c(1,2,3.3))
```

Then you feed this, as well as the `lm` object, into `predict()`

```
predict(object = lm_hibbs, newdata = newdf)
```

```
##           1           2           3
## 49.30818 52.36870 56.34739
```